

SHIVAM SHARMA

shivamsharma2023@gmail.com | (585) 481-9927 | [linkedin.com/in/shivamsharma-ai](https://www.linkedin.com/in/shivamsharma-ai) | shivamsportfolio.webflow.io

April 28, 2026

Hiring Team

Agoda — ML & LLM Platforms

Bangkok, Thailand

Dear Agoda Hiring Team,

I am applying for the Technical Product Manager, LLM/ML Domain role. The scope you describe — defining and evolving the product vision and roadmap for ML & LLM Ops platforms, owning model lifecycle and evaluation, supporting RAG and prompt management, balancing experimentation with governance and cost efficiency, and leveling the platform up for AI coding agents — is the architecture I just operationalized end-to-end at smaller scale at ESC Partners.

At ESC I built **Maya AI**, the agentic LLM platform behind a **200K+** resident utility on Oracle CCS. Maya unifies voice, chat, SMS, email, and mobile and autonomously executes billing, payment-plan, start/stop service, and outage workflows; it handles **70–80%** of first-touch volume across every channel. The parts most relevant to your role: an **LLM Ops platform** on AWS Bedrock that routes between Claude Opus, Sonnet, and Nova Pro against per-intent cost / latency / adoption SLOs; a **17-admin-API** governance layer with **24-hour key rotation** and tenant-level audit trail; deterministic HITL fallbacks via a one-click takeover cockpit; and a real-time observability surface with sentiment, confidence, and conversation history. I shipped it in **2.5 months** as the sole AI hire and the platform is now rolling out at Moreno Valley Electrical Utility and piloting at Delta Utilities.

Adjacent to Maya, I shipped **VaultIQ**, an air-gapped Claude Opus + Postgres/pgvector intelligence platform that runs entirely inside the customer VPC, with Terraform-defined per-customer deploys and role-based access. It is the **RAG-pipeline ownership** pattern Agoda's platform needs to support at scale: scoped knowledge corpus, retrieval-quality evaluation, freshness loops, least-privilege access. On Senior Assisted Living Locators I cut API spend **~70%** on a 300+ weekly consultation product through two-phase pre-filtering — the same caching / model-tier-down thinking that keeps LLM Ops cost efficient at scale.

Two pieces of context for the AI-native development direction Agoda is driving: I built Maya AI **directly on Claude Code**, and I built and use a **career-ops job-search pipeline** on Claude Code daily. First-hand, in-production experience with AI coding agents and the platform abstractions that make them reliable. Two prior exits (\$575K combined, \$5M+ deals influenced) taught me to ship in weeks and to write the governance and rollback plan before scaling tenants. I am open to relocating to Bangkok and to operating inside Agoda's AI transformation at scale.

You can see Maya behind a [live SALL demo](#), a [walkthrough video](#), and the rest of the work at shivamsportfolio.webflow.io. I would value the chance to discuss the ML / LLM Ops roadmap with you.

Sincerely,

Shivam Sharma