

SHIVAM SHARMA

shivamsharma2023@gmail.com | (585) 481-9927 | [linkedin.com/in/shivamsharma-ai](https://www.linkedin.com/in/shivamsharma-ai) | shivamsportfolio.webflow.io

April 27, 2026

Hiring Team

Capgemini — Data & AI

New York, NY

Dear Capgemini Hiring Team,

I am applying for the AI Product Manager role (Ref. **463746**). The scope you describe — vision and roadmap for agentic AI systems, an LLM evaluation framework over foundational and open-source models, RAG pipeline ownership, observability and incident-management integrations, and AI safety with deterministic fallbacks — is the architecture I just operationalized end-to-end at smaller scale.

At ESC Partners I built **Maya AI**, the agentic AI operating layer behind a **200K+** resident utility on Oracle CCS. Maya unifies voice, chat, SMS, email, and mobile, and autonomously executes billing, payment-plan, start/stop service, and outage workflows; it handles **70–80%** of first-touch volume across every channel. The parts most relevant to your role: an LLM evaluation framework on AWS Bedrock that routes between Claude Opus, Sonnet, and Nova Pro against per-intent cost, latency, and CSAT SLOs; a **17-admin-API** governance layer with **24-hour key rotation** and tenant-level audit trail; deterministic HITL fallbacks via a one-click takeover cockpit so non-technical supervisors retain control; and a real-time observability surface with sentiment, confidence, and conversation history for post-incident review. I shipped it in **2.5 months** as the sole AI hire and the platform is now rolling out at Moreno Valley Electrical Utility and piloting at Delta Utilities.

Adjacent to Maya, I shipped **VaultIQ**, an air-gapped Claude Opus + Postgres/pgvector intelligence platform that runs entirely inside the customer VPC, with Terraform-defined per-customer deploys and role-based access. It is the same RAG-pipeline ownership pattern Capgemini wants to apply to internal AIOps: scoped knowledge corpus, retrieval-quality evaluation, freshness loops, and least-privilege access. On Senior Assisted Living Locators I cut API spend **~70%** on a 300+ weekly consultation product through two-phase pre-filtering — the same cost / latency / model-routing thinking I would bring to Capgemini's production-engineering and incident-management surfaces.

Two exits before this (\$575K combined, \$5M+ deals influenced) taught me to ship in weeks, not quarters, and to write the governance and rollback plan before scaling tenants. I am ready to relocate to New York and to operate inside Capgemini's delivery scale.

You can see Maya behind a [live SALL demo](#), a [walkthrough video](#), and the rest of the work at shivamsportfolio.webflow.io. I would value the chance to discuss the agentic-AI roadmap with you.

Sincerely,

Shivam Sharma