

Application Q&A: Freed Product Manager

Pre-filled answers for the application form and likely screen-1 questions.

Why Freed?

Doctors are quitting in masses, and Freed is the team actually moving the needle on it. 20,000 paying clinicians, 70,000 patient notes daily, 2.5 million cumulative hours saved since 2023. That is a product I would want my own family's doctors to use.

The product motion is the same one I run at ESC: AI that gives a non-technical expert hours of their day back. Maya AI for utility supervisors does for energy what Freed's AI scribe does for clinicians: removes the administrative load so the expert can focus on the work that actually matters.

Sequoia, Daniel Gross, and ex-Meta founders mean the bar is high. That is the team I want to be on.

Why this role?

Four lines from the JD match how I already work:

1. **Own end-to-end product lifecycle.** Maya AI: vision to launch in 2.5 months as sole AI hire.
2. **Make clear, well-reasoned prioritization decisions.** Founder track at CloudApproach + Approachables required this constantly across two simultaneous ventures.
3. **Translate clinician needs into focused, practical product solutions.** ESC: translated NEP utility supervisors' Oracle CCS pain into a 5-channel autonomous agent platform.
4. **Use AI tools to accelerate your own workflows.** Daily Claude Code operator. PRDs as Claude conversations first. Maya AI was prototyped in Claude before any code shipped.

Why are you leaving your current role?

Scope and stage. ESC has been a great place to build Maya AI, but my next role should put me inside a team where the work I have been shipping for one customer at a time becomes the product at platform scale.

Compensation expectations?

Open to the posted range. My priority is total package and equity context given the stage and the scope. If we are aligned on the work, we can find a structure.

Tell me about a time you took ownership and delivered results.

Maya AI is the cleanest example. ESC had no AI track when I joined. The CRO wanted Oracle CCS to be smarter, faster, and multi-channel. There was no design doc, no eval framework, no team.

I started with two weeks of customer discovery at NEP. Sat with supervisors. Watched them dig through Oracle CCS to answer a single resident question. Mapped the workflow. Found the failure points: language barriers, response time, manual data entry, no escalation path.

I wrote a one-page product narrative. Picked the stack: AWS Bedrock for model routing, n8n for workflow orchestration, FastAPI for the integration layer, React for the supervisor cockpit, Microsoft Teams adaptive cards for human approval, 11 Labs for 26-language voice.

I built it on Claude Code, end-to-end, in 2.5 months. Single AI hire. Maya AI is live at NEP with 200K+ residents, rolling out at Moreno Valley Electrical, piloting at Delta Utilities. It handles 70 to 80% of first-touch support autonomously and removes 100+ manager emails per day. Projected savings at NEP are around \$225K per month.

Founder mode is how I worked then. It is how I work now.

What AI tools are you currently using today and how are you using them?

- **Claude Code:** my daily driver. Maya AI's supervisor cockpit, email-triage agent, React console, and FastAPI handlers all started as Claude Code sessions before any code shipped. PRDs are usually Claude conversations first, written documents second.
- **AWS Bedrock:** production routing layer for Maya AI. I route across Claude Opus (hard agent-of-record decisions), Claude Sonnet (default conversational tier), and Nova (cheap classifier-style intent extraction). Confidence distributions and cost-per-call drive the routing.
- **n8n:** orchestrator for three production agentic workflows. Conference Outreach (AI drafts, Teams approval cards, Outlook auto-send), Maya AI's email-triage routing, and the supervisor approval flow with adaptive cards.
- **11 Labs:** voice surface for Maya AI's 26-language coverage.
- **Claude (web app), ChatGPT, Perplexity:** mixed deliberately for thinking work. Claude for brainstorming, ChatGPT for second-opinion checks, Perplexity for fact-finding with citations.
- **Cursor and GitHub Copilot:** in-editor pair-programming inside long-lived codebases.

What changed when I started shipping with these tools daily: I went from waiting on engineering capacity to shipping most product ideas in hours instead of weeks. The 2.5-month Maya AI build is the proof.

What is your biggest accomplishment?

Career-wise: Bootstrapping and selling CloudApproach + Approachables for \$575K combined. Two ventures across UK and MENA, 200+ agency clients, 80+ real-estate partners, \$5M+ in deals influenced, 95% client retention post-acquisition. Two exits before I turned 25.

Last six months: Maya AI live at NEP with 200K+ residents in 2.5 months as the sole AI on staff. The team had never shipped AI before. They have shipped it now.

Are you authorized to work in the United States?

I am on STEM OPT (F-1 visa) through 2027 and am H-1B eligible. Happy to walk through the timeline in our first conversation.

When can you start?

Immediately. Two-week handoff at ESC for the Maya AI roadmap and supervisor cockpit. Day-one work in parallel.

Open to relocation?

Yes. Open to San Francisco within 30 days of offer.

Anything else you want us to know?

Three proof links if you want to see the work:

- SALL live: <https://ai-sales-assistant-frontend.onrender.com/>
- SALL walkthrough: https://www.youtube.com/watch?v=eCkP7_ZI348
- Portfolio: <https://shivamsportfolio.webflow.io/>

Happy to walk through a code-level Maya AI dive at any point in the loop.