

LinkedIn Messages: Morgan Stanley, AI Product Manager

Morgan Stanley does not name a hiring contact on this requisition, and no current Head of Firmwide AI has been publicly announced. The two people below are verified, correctly titled, and sit directly over the work this seat supports. Madhu Coimbatore is the sharper first move because the FINOS AI governance angle connects straight to the governance-measurement scope of the req.

A referral matters more than the application here. The requisition has been open since May 14, 2026 with 200+ applicants.

Madhu Coimbatore, Managing Director, Head of Firmwide AI Platform

LinkedIn: <https://www.linkedin.com/in/mcoimbatore/> Verified anchor: quoted in the FINOS press release launching the FINOS AI Fund on 23 June 2026, where Morgan Stanley is a Founding Premier Member. He framed the need for a shared baseline for how AI governance is understood and applied, and said Morgan Stanley actively uses the FINOS AI Governance Framework in new vendor onboarding.

Short (connection request, ~295 chars)

Hi Madhu, just applied to the AI Product Manager role supporting GenAI reporting and insights. Your FINOS point about needing a shared baseline for AI governance is the exact problem I hit at smaller scale: I built the measurement layer over a production agent platform for a 200K-resident utility. Would love to connect.

Long (InMail or post-connection)

Hi Madhu,

Just applied to the AI Product Manager role in the Technology COO org, the one covering GenAI adoption, usage, governance, and business impact reporting. Reaching out because your comment in the FINOS AI Fund announcement, about needing a shared baseline for how AI governance is understood and applied, is the same problem I ran into on a much smaller stage.

I am the first AI hire at ESC Partners. I built Maya AI end to end in 2.5 months: an autonomous agent platform on Oracle CCS now live for a 200K-resident utility, with two more in rollout. The part relevant to your world is what came after launch. Nobody could tell whether it was working, so I built the supervisor cockpit: confidence scoring on every response, live sentiment, and human takeover recorded per conversation. Then I set containment, confidence, and takeover frequency as the reported measures of quality, tracked per intent rather than blended, because blended numbers hide the intents that are actually failing.

What I will be upfront about: I have never worked inside an organization of Morgan Stanley's size, and I have not used Tableau or Power BI. I built the reporting surfaces myself in React over FastAPI, defined the metrics, and defended those definitions to the people being measured by them.

If you have fifteen minutes in the next couple of weeks, I would value your read on how the Firm is thinking about measuring GenAI now that SR 26-2 has left generative and agentic systems outside model risk scope. No agenda beyond that.

Shivam

H. David Wu, Managing Director, Firmwide AI Head of Product and Architecture Strategy

LinkedIn: <https://www.linkedin.com/in/hdavidwu> Verified anchor: spoke at the Global WealthTech Summit USA as Managing Director, Head of Knowledge Management and Generative AI. Separately quoted in OpenAI's Morgan Stanley case study on the Firm's evaluation discipline and on zero data retention being the contractual unlock for internal trust.

Short (connection request, ~290 chars)

Hi David, just applied to the AI Product Manager role covering GenAI reporting and insights. Your point in the OpenAI case study about evaluation discipline is what I spent the last six months on: confidence, override, and takeover tracked per intent on a live agent platform. Would love to connect.

Long (InMail or post-connection)

Hi David,

Just applied to the AI Product Manager role in the Technology COO org, covering GenAI adoption, usage, governance, and business impact reporting. Reaching out because of what you said in the OpenAI case study about evaluation discipline, and about zero data retention being the contractual unlock for internal trust. Both of those are the practitioner's version of this problem rather than the slideware version, which is why I wanted to write to you directly.

I am the first AI hire at ESC Partners, where I built Maya AI end to end in 2.5 months and then built the thing that measures it. Every conversation carries a confidence score, a sentiment reading, and a takeover record, and that instrumentation is what let us say with evidence that 70 to 80% of first-touch volume was handled autonomously rather than just believing it. Approvals route through Microsoft Teams adaptive cards so a named person signs off before the agent acts.

The trust point lands for me because of VaultIQ, an air-gapped financial intelligence platform I built that runs Claude Opus on Bedrock entirely inside the customer VPC, with SSO and role-based access deciding which analysts see which slices. I have never had the luxury of designing an AI system first and adding the controls afterward.

I will be straight about the gaps: no time inside a firm of Morgan Stanley's scale, and no Tableau or Power BI. I built the dashboards myself and defined what they measured.

If you have fifteen minutes, I would love to hear how the Firm is approaching the jump from measuring one product well to measuring a portfolio. Happy to work around your calendar.

Shivam

Notes

- Lead with Madhu. The FINOS governance angle is specific, recent, and directly adjacent to this req's scope.
- David Wu is the better second message and the stronger product conversation if Madhu does not respond within a week.
- Michael A. Pizzi (EVP, Global Head of Technology and Operations, ranked #27 on American Banker's Most Innovative People in Finance 2026 for DevGen.AI) is too senior to approach cold. Reference the organization, not the person.
- Do not mention Jeff McMillan as current Head of AI. He has left the Firm. Do not name any successor; no public announcement exists.
- Do not raise visa status in outreach. Morgan Stanley handles it as an application-form question. Save it for the recruiter screen.
- Name the gaps once in the long version and move on. This audience will find them anyway, and volunteering them reads as calibration rather than weakness.