

Dear Vivek, Pratyush, and the Sarvam team,

Just applied to the Product Manager (Models) role. Sending this because the JD describes exactly the model work I have been shipping at smaller scale at ESC, and I want to make sure the application lands.

I am the first AI hire at ESC Partners. I built **Maya AI** from scratch in **2.5 months** on Claude Code: an autonomous, event-driven agent platform on Oracle CCS for a **200K-resident utility**. The voice surface runs on **11 Labs with 26-language coverage**; the routing layer evaluates AWS Bedrock model selection across Claude Opus, Sonnet, and Nova; the eval framework tracks confidence distributions, sentiment shifts, and takeover frequency as service-level objectives. The supervisor cockpit gives non-technical operators a one-click override on every uncertain action. The product motion (own model lifecycle, set quality bar, ship and monitor) is the same one the Sarvam Model PM seat names.

Where I am not literal: I have not authored in-house ASR or TTS training pipelines. I have shipped production speech surfaces through 11 Labs and evaluated their quality (latency, comprehension, takeover rate) as part of Maya AI's deployment loop. My ML-eval rigor comes from Simon Vision Consulting on the City of Rochester fraud-detection engagement, where I lifted XGBoost accuracy from **62% to 73%** with SHAP-attributed explainability. Model evaluation, quality metrics, and human-readable explanations were the deliverable, not the by-product. The WER / CER / MOS framework is a one-week ramp from where I sit.

What draws me to Sarvam specifically: sovereign Indian AI is the most important infrastructure bet on the subcontinent for the next decade. Sarvam-105B (Indus), the Sarvam Kaze hardware launch, and the NVIDIA / Accel / HCLTech round all signal a company that ships at unicorn velocity. The Model PM seat sits at the layer where the decisions actually compound.

Before ESC, I bootstrapped and exited two ventures across UK and MENA (**\$575K combined exit, 95% post-acquisition retention**, \$5M+ in deals influenced on a \$200K marketing budget). Founder mode is how I have always operated.

One open question I would love to discuss in screen one: the JD lists Bengaluru on-site. I am US-based on STEM OPT through 2027. If the seat is hard-Bengaluru, I would relocate for the right role; if there is any remote-from-US flexibility for the first year, I would love to know. Either way, the substance match is real and I would rather have the conversation early.

SALL live: ai-sales-assistant-frontend.onrender.com

SALL walkthrough: youtube.com/watch?v=eCkP7_ZI348

Portfolio: shivamsportfolio.webflow.io

Shivam Sharma

shivamsharma2023@gmail.com | (585) 481-9927 | linkedin.com/in/shivamsharma-ai | shivamsportfolio.webflow.io