

Dear Vivek, Pratyush, and the Sarvam team,

Just applied to the Product Manager (On-Device and Edge AI) role. Sending this because the JD describes enterprise AI inside the customer security perimeter, which is exactly the surface I have been shipping at ESC for the last five months.

I am the first AI hire at ESC Partners. **VaultIQ** is the closest analog I can offer to Sarvam's on-device thesis: an air-gapped Claude Opus financial intelligence platform running **inside the customer VPC** with no data egress, multi-tenant access controls, SSO, role-based permissions on Postgres with pgvector, Terraform-deployed for repeatable customer-perimeter installs. Same Claude Code build pattern as Maya AI; the difference is the deployment perimeter. The data-residency, sovereign-cloud, and on-premise reflexes the JD calls out are the muscle memory I built shipping VaultIQ.

Maya AI is the multi-tenant SaaS side of the same surface: an autonomous, event-driven agent on Oracle CCS for a **200K-resident utility**, with DynamoDB tenant governance, role-based access, **17 admin APIs**, **24-hour key rotation**, and deep Microsoft ecosystem integration (Teams adaptive cards for HITL approval, Outlook for the email-triage agent). The supervisor cockpit gives non-technical operators one-click human takeover on every uncertain action; agent auditability is built into the eval framework.

Where I am honest about the gap: I have not directly shipped MDM or Intune fleet management. The transferable rigor is multi-tenant SaaS, OAuth2 SSO, role-based access controls, Microsoft Teams and Outlook integration, and a working understanding of how IT admins evaluate enterprise AI tools (gained directly from selling and supporting 200+ enterprise agency clients at CloudApproach across UK and MENA, where IT admin sign-off was the gating step for every contract). Entra ID, SAML, and SCIM follow naturally from the OAuth2 and SSO patterns I already ship. The first 30 days at Sarvam would close the MDM-specific gap.

On the structured-experiment side: I shadow-channel A/B every new Bedrock model against Maya AI's takeover-frequency SLO for two weeks before any production rollout. The same discipline applies to enterprise fleet deployments: phased rollout from NEP to Moreno Valley Electrical to Delta Utilities, each phase gated by the eval framework before the next one starts.

Before ESC, I bootstrapped and exited two B2B ventures across UK and MENA (**\$575K combined exit**, **95% post-acquisition retention**, **\$5M+** in enterprise deals influenced on a **\$200K** marketing budget). Direct revenue accountability with enterprise buyers is how I have always operated.

One open question I would love to discuss in screen one: the JD lists Bengaluru on-site. I am US-based on STEM OPT through 2027 and willing to relocate. If there is any remote-from-US flexibility for the first 6 to 12 months while the seat ramps (especially relevant for an on-device line working with silicon partners in the US), I would love to know. Either way, the substance match is real and I would rather have the conversation early.

SALL live: ai-sales-assistant-frontend.onrender.com

SALL walkthrough: youtube.com/watch?v=eCkP7_ZI348

Portfolio: shivamsportfolio.webflow.io

Shivam Sharma

shivamsharma2023@gmail.com | (585) 481-9927 | linkedin.com/in/shivamsharma-ai | shivamsportfolio.webflow.io